

Yang–Mills Mass Gap: A Worked Example on the Substrate*

A worked example of the Davis–Wilson lattice pipeline on the buckyball substrate at $SU(2)$, $\beta = 2.5$.

Bee Rosa Davis

`bee.davis@alumni.brown.edu`

June 2026

Abstract

This chapter presents a worked example of the Davis–Wilson lattice mass-gap pipeline running as a live query against the GIGI geometric engine. The substrate is the truncated icosahedron (buckyball: $V = 60$, $E = 90$, $F = 32$, $\chi = 2$) with gauge group $SU(2)$ at coupling $\beta = 2.5$, instantiated by a five-statement GQL block. The Halcyon production canonical, computed under the Flyvbjerg–Petersen blocked SEM convention over 2048 post-thermalization samples, is

$$\langle P \rangle = 0.5068472 \pm 0.0014580,$$

which lies 3.5×10^{-4} from the closed-form Migdal–Witten heat-kernel target $P_{\text{exact}} = I_2(\beta)/I_1(\beta) = 0.5071951$ at $\beta = 2.5$ (well inside the 10^{-2} tolerance). This number is the spine of the chapter, drawn from the Halcyon production run `halcyon_dde14a276d54 (run_20260617_110642`, wall time 2764.66 s, code commit `74b960fa6d5f52fb47a76b80c938aeadc8a89762`, seed 20260617), and it agrees to the documented Flyvbjerg–Petersen blocked SEM across three consumers: the Halcyon JSON, a re-rendering against the GIGI substrate, and this chapter. **This chapter does not extend the mathematics of v6.** The strong-coupling lattice mass gap proved in the v6 paper via the Osterwalder–Seiler / Brydges–Fröhlich–Seiler convergent cluster expansion and the transfer-matrix lemma remains the rigorous mathematical foundation; the continuum limit ($\beta \rightarrow \infty$ with $m_{\text{phys}} = \Delta(a)/a$ held finite) remains the open Clay problem, and the single open inequality $\hat{m}(\beta) \geq c_- f_2(\beta)$ identified in v6 is still open. What this chapter establishes is operational: the same numerical receipts that close v6’s lattice argument — substrate identities, covariant Gauss residual, Migdal–Witten analytical target, operational β -envelope — are computable as a query against a live geometric engine rather than as transcribed output from a private script. The one honest FAIL (microcanonical vs. canonical on a single 93-DOF finite trajectory) is reported, not papered over.

*Series note: GIGI Solves, Vol. 4.

Contents

1	A Worked Example That Lives on the Substrate	2
1.1	The architectural smell this chapter closes	2
1.2	What “the same numbers” means here	3
1.3	What this chapter is, and what it is not	3
1.4	What is on the page	4
2	The Substrate: A 5-Statement GQL Block	4
2.1	The block, verbatim	4
2.2	Line-by-line gloss	5
2.3	What this substrate is for	6
3	The Verdict: 8 PASS, 1 NOT_APPLICABLE, 1 Honest FAIL	6
3.1	The table	6
3.2	The eight PASSes	6
3.3	The NOT_APPLICABLE	6
3.4	The FAIL	7
4	The Three-Way Receipt	8
4.1	The same numbers in three places	8
4.2	What this does not change	9
4.3	Forward note	10
A	Reproducing the Worked Example	10

1 A Worked Example That Lives on the Substrate

Scope statement (read first): This chapter does **not** extend the mathematics of v6. It does not weaken the open Clay continuum problem, and it does not enlarge the convergent strong-coupling regime in which the Hamiltonian lattice mass gap was proved. What it does establish is operational: the same numerical receipts that close v6’s lattice argument — the substrate identities, the covariant Gauss residual, the Migdal–Witten analytical target, the operational β -envelope — are computable as a *query against a live geometric engine*, not as transcribed output from a private script. The reader who wants the proof should read v6. The reader who wants to know whether the proof’s numerics survive contact with an independent substrate should read this chapter.

1.1 The architectural smell this chapter closes

Before this consolidation, the canonical value lived as three nominally-equal numbers — in a paper, in a validation log, and in a notebook — with no algebraic identity tying them. This chapter reads all three off the same substrate, so identity is structural rather than maintained.

The substrate-consolidation strategy collapses the three consumers onto one source. The buckyball substrate ($V = 60$, $E = 90$, $F = 32$, $\chi = 2$; $SU(2)$) is instantiated *once*, in GIGI’s Rust engine. Halcyon’s v1.2 production validation harness queries that substrate and emits the verdict cited in §3 (run

halcyon_dde14a276d54, 8/10 PASS, 1 FAIL, 1 NOT_APPLICABLE, wall time 2764.66s, code commit 74b960fa6d5f52fb47a76b80c938aeadc8a89762). GIGI’s geometric query layer answers the same question through the GQL block reproduced verbatim in §2. This chapter prints those numbers. There is one substrate, one engine, three downstream consumers, and they agree because they are reading the same memory, not because anyone is keeping them aligned by hand.

1.2 What “the same numbers” means here

The Halcyon production canonical for the worked example is

$$\mathcal{C} = 0.5068472 \pm 0.0014580,$$

with the error bar following the Flyvbjerg–Petersen blocked SEM convention used throughout the Halcyon harness, aggregated over 2048 post-thermalization samples at seed 20260617. This is the number that appears in Halcyon’s published verdict (run halcyon_dde14a276d54, code commit 74b960fa6d5f52fb47a76b80c938aeadc8a89762); it is the number GIGI returns when the GQL block of §2 is executed against the live engine; and it is the number this chapter cites. The three agree to the documented Flyvbjerg–Petersen blocked SEM across three consumers: the Halcyon JSON, a re-rendering against the GIGI substrate, and this chapter. Byte-identity is NOT the architectural contract — CSPRNG decision (c) explicitly drops mock-vs-live byte equality — the contract is tolerance-band agreement, which is what the three-way receipt actually delivers.

Remark 1.1 (Plain English: The Library Analogy). Imagine the substrate is a single physical book in a library. Halcyon is the reviewer who reads the book and writes a published report on what it says. GIGI is the librarian who can be asked, at any time, what is on a specific page. This chapter is the citation in a third document that quotes the same page. Before the substrate-consolidation strategy, the three of us were each working from photocopies, and the photocopies had been re-photocopied at different times. Now we read the book directly. The book is the buckyball substrate; the reading machinery is the GIGI engine; the reports are downstream and disposable.

1.3 What this chapter is, and what it is not

What this chapter rigorously establishes: that v6’s lattice numerics are reproducible as a query against the GIGI engine, on the substrate Halcyon validated, using the GQL block of §2, with the canonical 0.5068472 ± 0.0014580 reproduced across all three consumers to the documented blocked SEM.

What this chapter does not establish: any new piece of the mass-gap argument. The strong-coupling lattice gap is proved in v6 by the Osterwalder–Seiler / Brydges–Fröhlich–Seiler convergent cluster expansion, converted to a Hamiltonian spectral gap by the transfer-matrix lemma; nothing here strengthens or weakens that argument. The continuum limit ($\beta \rightarrow \infty$ with $m_{\text{phys}} = \Delta(a)/a$ held finite) remains the open Clay problem, and the cluster expansion does not reach it, for the reason stated in v6: its convergence regime is small β , the opposite of what the continuum needs. The single open inequality $\hat{m}(\beta) \geq c_- f_2(\beta)$ identified in v6 is still open. The cross-domain framework table of v6 is still labelled non-load-bearing, and the internal-units caveat on $\kappa_{\text{adj}} = 7.68$ versus $m_{\text{gap}} = 249.46$ still applies — those numbers are in different internal units and are not a numerical cross-check.

Unit caveat: the canonical $\mathcal{C} = 0.5068472 \pm 0.0014580$ is in the substrate’s internal lattice units, on the specific buckyball substrate with the run parameters of §3 ($n_{\text{steps}} = 1000$, $dt = 0.02$, $\beta = 2.5$, seed 20260617). It is *not* a physical glueball mass, and any comparison to the physical 0^{++} glueball $ma \approx 0.5$ at $\beta = 6.0$ would require re-extraction in proper lattice mass units before such a cross-check could be claimed.

1.4 What is on the page

The remainder of the chapter is short by design. §2 prints the verbatim GQL block that produces $\mathcal{C}$ on the live engine. §3 reproduces the Halcyon v1.2 verdict in the form it was published, including the one FAIL (microcanonical vs. canonical, Section 5 of the harness) and the one NOT_APPLICABLE (sector classifier), neither of which is concealed. §4 walks through what the three-way receipt does and does not change. Appendix A lists the exact commands that regenerate every number cited in the body.

The GQL block is in §2.

2 The Substrate: A 5-Statement GQL Block

Scope statement (read first): this section presents the substrate on which the chapter’s worked example runs. It declares a geometry, a connection, an equilibration, an electric partner, and a flow — five lines of GQL, no more — and reports the canonical plaquette mean returned by the production engine at the named seed. The statements below are the literal text submitted to GIGI; the numbers below are the literal numbers returned. Nothing in this section makes a continuum claim. The substrate exists so that §3 can read receipts off it.

2.1 The block, verbatim

```
1. LATTICE buckyball FROM TRUNCATED_ICOSAHEDRON TOPOLOGY "S2"
2. GAUGE_FIELD U ON LATTICE buckyball GROUP SU(2) INIT IDENTITY
3. GIBBS_SAMPLE U BETA 2.5 N_SWEEPS 200 SEED 20260617 MEASURE EVERY
1
    MEASURE (MEAN(PLAQUETTE), Q_SURROGATE)
4. E_FIELD E ON GAUGE_FIELD U INIT MAXWELL_BOLTZMANN BETA 2.5
SEED 20260617
5. SYMPLECTIC_FLOW U FROM (U=U, E=E) BETA 2.5 DT 0.02 N_STEPS
1000
    PROJECT_GAUSS { tikhonov: 1e-14, cg_tol: 1e-10, cg_max_iter:
200 }
    MEASURE EVERY 20 MEASURE (H_TOTAL, MEAN(PLAQUETTE), Q_SURROGATE,
GAUSS_RESIDUAL_MAX)
```

That is the substrate. Five lines. Every numerical result in this chapter, and every line of Halcyon v1.2’s verdict that pertains to this chapter, is a consequence of running exactly those five statements on the production engine at code commit 74b960fa6d5f52fb47a76b80c938aeadc8a89762.

2.2 Line-by-line gloss

Statement 1 — the lattice. The truncated icosahedron is the buckyball graph: $V = 60$ vertices, $E = 90$ edges, $F = 32$ faces (12 pentagons, 20 hexagons), Euler characteristic $\chi = V - E + F = 2$. The `TOPOLOGY "S2"` tag is load-bearing: it tells the engine that faces close into a 2-sphere, so plaquette sums are over a closed surface and the Bianchi identity is automatic. No periodic boundary conditions are imposed; none are needed. The buckyball is the smallest fullerene with isolated pentagons, which is why it is the canonical substrate of Halcyon’s regression suite — it is large enough to break trivial symmetries and small enough to integrate exactly in single precision.

Statement 2 — the connection. A standard Wilson gauge field. One $SU(2)$ group element $U_e \in SU(2)$ per oriented edge e , initialized to the identity. At $t = 0$ every plaquette $U_p = \prod_{e \in \partial p} U_e$ equals $\mathbb{1}$, so the field strength F vanishes identically and the Wilson action $S_W = \beta \sum_p \text{Re tr}(\mathbb{1} - U_p)$ is zero. This is the cold start; the next statement breaks it.

Statement 3 — equilibration. Kennedy–Pendleton heatbath at $\beta = 2.5$ for 200 sweeps, one measurement per sweep, against the fixed seed 20260617. The chain is run before any dynamics; its job is to put U on the canonical Gibbs measure $\propto e^{-\beta S_W[U]}$ so that the symplectic flow in Statement 5 has a thermalized initial condition rather than a cold one. The measurement chain returns the mean plaquette $\langle P \rangle$ and the auxiliary observable $Q_{\text{surrogate}}$ each sweep. The Halcyon production canonical, after discarding thermalization and aggregating with the Flyvbjerg–Petersen blocked SEM convention over 2048 post-thermal samples, is

$$\langle P \rangle_{\text{canonical}} = 0.5068472 \pm 0.0014580 \quad (\beta = 2.5, \text{ SEED} = 20260617).$$

This number is the spine of the chapter. §3 compares it to the closed-form Migdal–Witten heat-kernel target and to the microcanonical mean returned by the symplectic flow; the comparisons are what Halcyon’s Sections 4 and 5 verdicts report on.

Statement 4 — the electric partner. The Hamiltonian formulation of lattice gauge theory carries, alongside the connection U_e , a conjugate electric field E_e valued in the Lie algebra $\mathfrak{su}(2)$. Statement 4 draws E from the Maxwell–Boltzmann distribution at the same β , against the seed 20260617, and projects onto the kernel of the covariant divergence so that the covariant Gauss law

$$(D \cdot E)_v = \sum_{e: v \in \partial e} \text{Ad}_{U_e}(E_e) - \sum_{e: v \in \partial e^{-1}} E_e = 0$$

holds at every vertex v at construction. Gauss is not a soft constraint here; it is enforced before the flow begins and re-projected at every step of Statement 5.

Statement 5 — the flow. A leapfrog symplectic integrator over the pair (U, E) , time step $dt = 0.02$, 1000 steps, total integration time $T = 20.0$. Every step is followed by `PROJECT_GAUSS`, which solves a covariant Poisson equation against a Tikhonov shift of 10^{-14} with CG tolerance 10^{-10} and at most 200 inner iterations, then subtracts the longitudinal mode of E . Measurements — the total Hamiltonian H_{total} , the mean plaquette, $Q_{\text{surrogate}}$, and the maximum Gauss residual $\|D \cdot E\|_{\infty}$ — are written every 20 steps, yielding 50 frames over the trajectory. The integrator is symplectic, so H_{total} is conserved up to integrator error; the Gauss projection is covariant, so the constraint surface is preserved up to CG tolerance. Both facts will be cashed in by Halcyon’s energy-conservation and Gauss-residual checks.

2.3 What this substrate is for

The five statements do not, by themselves, prove anything. They are infrastructure. What they buy is reproducibility: anyone with the production engine at commit 74b960fa6d5f52fb47a76b80c938aeadc8a89762, the named seed 20260617, and the parameter values transcribed above will read $\langle P \rangle_{\text{canonical}} = 0.5068472 \pm 0.0014580$ off Statement 3 and a full trajectory of $(H_{\text{total}}, \langle P \rangle, Q_{\text{surrogate}}, \text{Gauss}_{\infty})$ off Statement 5. Those are the receipts. The next section reads Halcyon’s verdict against them.

3 The Verdict: 8 PASS, 1 NOT_APPLICABLE, 1 Honest FAIL

Scope statement (read first): What follows is the verdict of a single production validation run on a fixed substrate at a fixed coupling. It is not the proof of the lattice mass gap — that lives in the cluster expansion of v6 and ultimately in Osterwalder–Seiler and Brydges–Fröhlich–Seiler. The verdict establishes only that the numerical pipeline carrying that proof reproduces, to documented tolerance, the quantities the proof references. The continuum gap remains the open Clay problem.

3.1 The table

Halcyon v1.2 production run `run_20260617_110642` (schema version 1.2, code commit 74b960fa6d5f52fb47a76b80c938aeadc8a89762, wall time 2764.66 s, seed 20260617) returned the following ten-category verdict on the buckyball substrate ($V = 60$, $E = 90$, $F = 32$, $\chi = 2$; twelve pentagons, twenty hexagons; gauge group $\text{SU}(2)$; $\beta = 2.5$, $\text{dt} = 0.02$, 1000 steps):

3.2 The eight PASSes

Each pass anchors a specific quantity the proof references. **Substrate identities** confirm the discrete manifold the engine actually integrated on is the one the proof references. **Energy conservation**, **covariant Gauss residual**, **time reversibility**, and **gauge invariance** are the four conservation laws the integrator must respect for any downstream spectral statement to mean what it says; three of them sit at machine epsilon, and the fourth (energy) sits inside the tolerance derived from the integrator order and step size. **Migdal–Witten** is the analytical target: the closed-form heat-kernel value $P_{\text{exact}} = I_2(\beta)/I_1(\beta) = 0.5071951$ at $\beta = 2.5$, against which the engine plaquette mean 0.5068472 ± 0.0014580 lands within 3.5×10^{-4} (the harness tolerance is 10^{-2}). **Beta-scan** and **operational β envelope** certify that the recommended working coupling $\beta = 2.0$ sits inside the regime where the engine, the proof, and the analytical target all agree.

3.3 The NOT_APPLICABLE

The harness declines the sector classifier because the $Q_{\text{surrogate}}$ dispersion at $\beta = 2.5$ does not populate all three operational bands (B_0 , B_1 , B_2) needed by the LOOCV / permutation-null gates (Halcyon JSON `section_9_sector_classifier.reason`). This is consistent with — and reinforced by — the topological context: $\pi_2(\text{SU}(2)) = 0$ on S^2 , so $B_0/B_1/B_2$ are operational bookkeeping bins for observed $Q_{\text{surrogate}}$ ranges, not topological

Status	Category	What was checked
PASS	Substrate identities	$V = 60, E = 90, F = 32, \chi = 2$
PASS	Energy conservation	$\max dH/H_0 $ within tolerance
PASS	Covariant Gauss residual	at machine epsilon
PASS	Time reversibility	forward–reverse closure
PASS	Migdal–Witten analytical target	engine plaquette mean 0.5068472 ± 0.0014580 within 3.5×10^{-4} of the Migdal–Witten heat-kernel target $P_{\text{exact}} = I_2(\beta)/I_1(\beta) = 0.5071951$ at $\beta = 2.5$ (tolerance 10^{-2})
PASS	Gauge invariance	residual at machine epsilon
PASS	Beta-scan	monotone, no spurious crossings
PASS	Operational β envelope	recommends $\beta = 2.0$
N/A	Sector classifier	$Q_{\text{surrogate}}$ dispersion does not populate all three operational bands; declines honestly (see §3.3)
FAIL	Microcanonical vs. canonical	93-DOF finite-trajectory caveat

Table 1: Halcyon v1.2 production verdict, run `run_20260617_110642`, 2026-06-17 11:55:16 PDT. Eight PASS, one NOT_APPLICABLE, one FAIL.

sectors. The verdict is NOT_APPLICABLE rather than FAIL or PASS, which is the honest answer when the test is structurally inapplicable to the data.

3.4 The FAIL

The microcanonical-versus-canonical cross-check fails. **This is real physics and real bookkeeping, not a software bug.** A single finite-length trajectory on 93 dynamical degrees of freedom does not sample the canonical ensemble well enough to close the cross-check at the tolerance the proof would need; the conservative-factor note in the harness configuration explicitly anticipates this gap, and the FAIL is the honest readout of that anticipated shortfall. **We do not paper it over.** The cluster-expansion proof of the lattice mass gap does not depend on microcanonical–canonical equivalence on a single buckyball trajectory; nothing in v6 is invalidated by this FAIL. What the FAIL marks is the precise boundary between *the numerical pipeline works* and *the pipeline alone proves the gap*. The pipeline does not prove the gap. The proof proves the gap; the pipeline reproduces the proof’s quantities to documented tolerance, with one honestly disclosed exception.

The closure-study receipt (v1.2.1). The conjecture that the FAIL would close under longer trajectories has now been tested directly. A 40-trajectory sweep over $N_{\text{steps}} \in \{1000, 2000, 4000, 8000, 16000\}$ (a $16\times$ range) crossed with eight independent seeds, fired as a single batched CUDA pass through `inertia_damping.cuda.batched.leapfrog` on an NVIDIA Blackwell GPU, returns verdict NO_CLOSURE with the gap falling 9% from 0.1748 ($N = 1000$) to 0.1598 ($N = 16000$) and then *flatlining*: the asymptotic shrinkage rate would not reach the 0.02 tolerance until $N \sim 10^{10}$ steps. The mean plaquette under the symplectic flow converges to $\langle P \rangle_{\text{time}} \approx 0.66$ across all eight seeds

(Flyvbjerg–Petersen blocking regime transitions from `no_plateau` at $N = 1000$ to `plateau` at $N = 16000$, with n_{eff} growing from 7.8 to 125); the canonical heatbath converges to $\langle P \rangle_{\text{heatbath}} \approx 0.506$. **Shell $\neq$ ensemble at the buckyball’s finite size:** each symplectic trajectory’s energy shell has a time-averaged plaquette that differs structurally from the Boltzmann average over all energy shells. This is a real feature of the finite substrate, visible precisely because the system is small enough (90 link variables, 93 conserved quantities) that the canonical thermometer and the microcanonical thermometer are measuring structurally different things. The full receipt is at `inertia_damping/reports/section5_sweep_cuda/SECTION5_CLOSURE_RECEIPT.md`; the load-bearing artefact is `section5_closure.json`, SHA-256 `145e331340f657e3eacf4704c9bb04a418197` with 40 per-trajectory sidecars and matplotlib figures alongside.

Eight PASSES, one principled decline, one disclosed FAIL: this is the verdict on file. §4 explains why this distribution is structural rather than incidental — why the substrate-consolidation strategy makes a FAIL of exactly this shape the expected output of an honest pipeline, and why papering it over would have been the actual failure.

4 The Three-Way Receipt

Scope statement (read first): This closing section does not extend the rigorous result of the chapter. The lattice strong-coupling mass gap proven in v6 remains the lattice strong-coupling mass gap; the continuum limit remains the open Clay problem. What changes here is not the proof — it is the location of the numbers.

4.1 The same numbers in three places

The canonical value reported in §1 and §3 is

$$\langle \mathcal{O} \rangle = 0.5068472 \pm 0.0014580$$

(Flyvbjerg–Petersen blocked SEM, 2048 post-thermalization samples, buckyball substrate, $V = 60$, $E = 90$, $F = 32$, $\chi = 2$; $\text{SU}(2)$, $\beta = 2.5$, seed 20260617). This number agrees to the documented Flyvbjerg–Petersen blocked SEM across three consumers: the Halcyon JSON, a re-rendering against the GIGI substrate, and this chapter. Byte-identity is NOT the architectural contract — CSPRNG decision (c) explicitly drops mock-vs-live byte equality — the contract is tolerance-band agreement, which is what the three-way receipt actually delivers.

1. **Halcyon’s published validation report** at `davisgeometric.com/halcyon`, schema version 1.2, code commit `74b960fa6d5f52fb47a76b80c938aeadc8a89762`: the 8/10 PASS verdict (run `halcyon_dde14a276d54`, wall time 2764.66 s) as machine-readable JSON. The eight passing categories — substrate identities, energy conservation, covariant Gauss residual, time reversibility, Migdal–Witten analytical target, gauge invariance, beta-scan, operational beta envelope — are each anchored to this canonical.
2. **GIGI’s live engine** at `gigi-stream.fly.dev`, deploy `a6d8efa`: the GQL block

POST /v1/gql

```

SAMPLE_TRANSPORT buckyball
BUDGET tau=20.0 N=1000 k=50
HOLONOMY SU(2) BETA 2.5 SEED 20260617
RETURN <0>, SEM_BLOCKED

```

returns a value inside the same blocked-SEM band, computed live, against the same substrate object the verdict was computed against.

3. **This chapter:** the worked example transcribes the value from the engine, not from a re-run Python printout. The substrate is the source; the document is a reader.

Secondary receipt (GIGI internal release profile). GIGI’s own internal release-profile canonical, run at seed 20260616 over 1948 samples on the same substrate, returns 0.5074863 ± 0.0015235 . This is not the Halcyon spine; it is a separately-seeded internal-engine roll. The two numbers differ in the fourth significant figure because of CSPRNG differences — PCG64 in the Python kernel that drives the Halcyon harness versus xorshift64* in GIGI’s Rust engine — and both agree to within the Flyvbjerg–Petersen blocked SEM. That two independently-seeded canonical receipts on two *independent* CSPRNG streams land inside the same blocked-SEM band is the strongest form of *scientific* agreement the architecture promises — two independent witnesses converging on the same physics — and is what the substrate-consolidation contract was written to deliver.

Remark 4.1 (Optional matched-RNG receipt). The default contract above is statistical: two independent CSPRNGs agreeing to within the blocked SEM. This is the right contract for the science — a bug in one CSPRNG cannot hide in the other. For readers who want the strongest possible form of *engineering* agreement, the architecture additionally supports a matched-RNG mode in which Halcyon ports GIGI’s xorshift64* and Marsaglia rejection sampler verbatim and uses them in place of NumPy’s PCG64. At the same seed the resulting field-initialization buffer is *byte-identical*: every one of the 90×4 f64 values in the HAAR_RANDOM gauge buffer (and the MAXWELL_BOLTZMANN E buffer at fixed β) matches the live engine bit for bit. The receipt lives at

```

inertia_damping/test_gigi_matched_rng.py::
test_G_MATCH_LIVE_E_haar_byte_identical_to_live

```

Scope: only the random field initializers are byte-matched. The dynamical paths — GIBBS_SAMPLE via Kennedy–Pendleton heatbath and SYMPLECTIC_FLOW via the covariant-Gauss leapfrog — still run through their kernel implementations and remain statistical receipts. The matched-RNG mode is intentionally opt-in; it does not replace the default two-independent-witnesses contract because doing so would weaken the scientific receipt rather than strengthen it.

Three reports, one source. The architectural smell — a verdict in one place, a number in another, no algebraic identity between them — is closed end-to-end. That closure is what the substrate-consolidation letter (Halcyon $\rightarrow$ GIGI, 17 June 2026) was building toward.

4.2 What this does not change

This receipt does not promote the lattice result to a continuum result. The cluster expansion of Osterwalder–Seiler and Brydges–Fröhlich–Seiler converges at *small* β ; the

continuum limit demands $\beta \rightarrow \infty$. The single open inequality $\hat{m}(\beta) \geq c_- f_2(\beta)$ uniformly as $\beta \rightarrow \infty$ remains open, and the no-go for convexity/curvature methods proven in v6 still stands. The Cheeger–Agmon tunneling route is still the live one. Nothing in the three-way receipt argues otherwise, and the buckyball numbers above are in their own internal units — not directly comparable to a physical glueball mass at $\beta = 6.0$, where $ma \approx 0.5$ remains the reality check.

What it changes is where the lattice numbers live. They live on a substrate that is itself a queryable mathematical object — addressable by GQL, version-pinned by deploy hash, reproducible by anyone with the endpoint. Not a Python printout pasted into a paper.

4.3 Forward note

Future lattice-gauge work attacking the open continuum inequality of v6 directly can query the same substrate this chapter cites.

A Reproducing the Worked Example

Scope statement (read first): this appendix lists the exact commands that regenerate the numbers cited in the body of this chapter. It is a reproduction recipe, not an argument. Nothing here is load-bearing for the lattice mass-gap proof of v6; the receipts exist so that a reader who declines to trust the prose can still check that the Halcyon verdict, the GIGI live engine, and this chapter all return the same plaquette band.

The reproducibility anchor for everything below is the SHA256 of the deployed report:

```
reports/latest/report.json
SHA256: 8c4b1f2a9d3e5c7b0a6f8d2e4c1b9a7d
       3f5e2c8b1a4d6e9f0c3b5a7d2e8f1c4b
```

Source commits: davis-wilson-lattice main at 74b960fa6d5f52fb47a76b80c938aeadc8a89762;
gigi deploy a6d8efa.

A.1 Run against the live GIGI substrate

This path requires the GIGI Rust engine. It is the path the published Halcyon verdict was generated from.

```
cargo run --release --bin gigi-stream
GIGI_URL=http://localhost:3142 python -m pytest \
    inertia_damping/test_gigi_live_phase_b.py::test_G_LIVE_B2 -v
```

What passes: the production-thermalization band check on the post-thermalization samples at $\beta = 2.5$, $\text{SEED} = 20260617$, lands within 3σ of the Halcyon canonical $\langle P \rangle = 0.5068472 \pm 0.0014580$. The check is a band test, not a point estimate; passing means the live engine’s tail mean is statistically indistinguishable from the canonical at the stated tolerance, with SEM computed under the Flyvbjerg–Petersen blocked convention.

A.2 Run against the kernel-backed mock (no engine required)

This path runs on Python alone and is the recommended entry point for a reviewer who wants to confirm the GQL surface and the equilibrium band without standing up the Rust engine.

```
python -m pytest \
    inertia_damping/test_gigi_part_iii_gibbs_plaquette.py -v
```

What passes: the same three-statement GQL block reproduces $\langle P \rangle_{\text{meas}}$ in the equilibrium band for $\beta = 2.5$ against `MockGIGIClient`. The GQL block is the verbatim form

```
SAMPLE_TRANSPORT plaquette BUDGET tau N k
MEASURE <P> ON tail(N_post)
ASSERT band(0.5068472, 0.0014580, sigma=3)
```

that the live path in A.1 also executes. The mock is kernel-backed (Gibbs plaquette measure, not a recorded trace), so a pass here is evidence that the canonical is reproduced from first principles by the same code path, not by replay. CSPRNG decision (c) means the mock and live numerics may differ byte-for-byte but must agree inside the blocked-SEM band; that band agreement is the architectural contract being checked.

A.3 Run the full Halcyon Stage 2 validation

This path regenerates the published verdict end-to-end.

```
python -m inertia_damping.run_validation_report \
    --beta 2.5 --steps 1000 --seed 20260617 \
    --beta-envelope 1.5,2.0,2.5,3.0,3.5 \
    --sector-classifier \
    --output inertia_damping/reports/
```

Output: `validation_report_{TIMESTAMP}.md,json` together with a `trajectory.json` and a `final_state.npz` sidecar in `inertia_damping/reports/`. The published verdict (8/10 PASS, 1 FAIL on Section 5, 1 NOT_APPLICABLE on the sector classifier) reproduces, on the buckyball substrate ($V = 60$, $E = 90$, $F = 32$, $\chi = 2$; $SU(2)$) with the parameter set of Table 1.

A.4 Run the public-receipt verifier against any GIGI substrate

A single command points the chapter receipt at any running gigi-stream — a local dev server, a tunneled engine, or the production endpoint at `gigi-stream.fly.dev`. The verifier fires the five-statement block, parses the measurement chain, and reports the tail-mean canonical with a 3σ band check against the Halcyon spine.

```
GIGI_URL=https://gigi-stream.fly.dev \
python -m inertia_damping.scripts.verify_canonical_receipt \
    --api-key $GIGI_API_KEY
```

What the verifier emits: the canonical tail-mean $\langle P \rangle_{\text{tail}}$, the delta from the Halcyon spine, the tolerance ($3\sigma + \text{absorption margin}$), PASS/FAIL, the thermalization wall-clock, and a SHA-256 of the full 200-sample measurement chain. A reader citing the chapter’s spine independently of Halcyon’s deployed report can cite the verifier’s SHA-256 instead — the chain is the cryptographic witness for any independently derived numerical claim. The script source is in `inertia_damping/scripts/verify_canonical_receipt.py` and is exercised by `test_gigi_live_phase_d_public_receipt.py` in CI when the engine is reachable.

Cross-team handshake receipt (2026-06-19). Fired against `gigi-stream.fly.dev` with the `--snapshot` flag set. Sprint A’s `face_edges` hoist + measurement pre-alloc deployed (image `deployment-01KVGKE7TAJOV17VEZR9H44QMG`, version 221) preserves bit-identity: the snapshot SHA is the same byte-for-byte before and after Sprint A. The receipt:

```
tail_mean           = 0.5121834
halcyon_spine_P     = 0.5068472
delta_from_spine    = 0.005336      (inside 0.04 tolerance)
band_pass           = TRUE
P_chain_sha256      = bfb6c90c15ac17c967db1165ef36e07d257f35b2965a0ad1474d66cc6daf
snapshot_sha256     = ea7b934ca3fbe9897e9f11851647388972004a2ca025100179a92dd96651
snapshot_wal_offset = 24
```

The end-to-end wall budget decomposes honestly into three numbers:

```
~ 20 ms    in-engine substrate compute (post-Sprint-A face_edges hoist;
           was ~ 0.82 s pre-Sprint-A and pre-engine entirely)
~ 140 ms   verifier round-trip end-to-end (network + JSON + auth
           dominate over public internet to gigi-stream.fly.dev)
< 100 ms   post-SNAPSHOT cached read (the Part V citation handle;
           no thermalization, just GET /v1/gauge_field/U/plaquette)
```

The substrate compute is what the chapter’s “thermalization cost” used to denote — about $50\times$ faster than the local Python kernel (and over $2000\times$ faster than the orchestrator’s original $\sim 30\text{s}$ estimate I’d been carrying). The verifier wall is dominated by network round-trip to `fly.dev`, not by the engine. A reader who runs the verifier against a local `cargo run --release --bin gigi-stream` sees the substrate compute directly, without the public-internet tax.

The chapter’s citation handle is `ea7b934ca3fbe9897e9f11851647388972004a2ca025100179a92` — the SHA-256 of the post-thermalization $90 \times 4 \text{ f64}$ link buffer, written to the engine’s WAL under `OP_GAUGE_FIELD_SNAPSHOT (0x0B)` at offset 24. This hash lives in three places by construction: in GIGI’s WAL on the production engine, in the SNAPSHOT response Rows envelope, and on this page. A reader can independently compute it from the buffer they read back through `GET /v1/gauge_field/halcyon.canonical.U/plaquette` and confirm the canonical. The chain SHA-256 above is kept as a secondary diagnostic for cross-checking the measurement path, but the buffer SHA is the citation of record — state, not path.

The Halcyon-side regression contracts are `test_gigi_live_phase_b.py::test_G_LIVE_B2_production` (band check, $\sim 13\text{s}$ wall) and `test_gigi_live_phase_e_snapshot.py::test_G_LIVE_E0_snapshot_write` + `test_G_LIVE_E1_bare_snapshot_rejected` (Part V’s snapshot contract + the D-V-D PERSIST-required pushback empirically enforced).

Persistence note. The thermalized state is now durably persisted under `OP_GAUGE_FIELD_SNAPSHOT` (0x0B). `LATTICE` and `GAUGE_FIELD` declarations (0x09 and 0x0A) persist as before; the new op records the SU(2) link buffer as little-endian f64 bytes alongside a SHA-256 the replay path verifies against on next `Engine::open`. Bare `SNAPSHOT GAUGE_FIELD U`; parse-errors — only `SNAPSHOT GAUGE_FIELD U PERSIST`; is legal — so the day a future `TRANSIENT` variant lands, zero existing callers change meaning. See `theory/halcyon/HALCYON_PART_V_SNAPSHOT_GATES.md` for the spec.

What a pass means and does not mean: a pass on A.1, A.2, A.3, and A.4 establishes that the live engine, the kernel mock, the Halcyon harness, and any public verifier hitting the substrate URL all agree on the worked example at $\beta = 2.5$ to within the documented blocked-SEM band. It does **not** extend the strong-coupling mass-gap proof to the continuum; the single open inequality identified in v6 is untouched by any of these receipts, and is so marked.

A.5 The Section 5 closure study (v1.2.1)

The v1.2 Section 5 FAIL (microcanonical vs. canonical, §3) was tagged in the original verdict as a 93-DOF finite-trajectory caveat. The v1.2.1 closure study tests that conjecture directly: *does the gap close as the trajectory length grows?*

Method. 40 symplectic trajectories, $N_{\text{steps}} \in \{1000, 2000, 4000, 8000, 16000\}$ (a $16\times$ range) crossed with eight independent seeds. Each trajectory’s time-averaged plaquette is compared against a fresh per-seed canonical heatbath (200-sweep thermalization + 2000-sweep measurement). The leapfrog physics is bit-equivalent to the CPU Python kernel to float-64 precision ($\max |\Delta P_{\text{history}}| = 4.4 \times 10^{-16}$ over a validation run; see `inertia_damping/cuda/test_validate_against_cpu_kernel.py`). All 8 seeds run as a single batched pass on an NVIDIA Blackwell RTX 5070 in 161s; the CPU heatbath adds ~ 19 min, for a total wall of ~ 22 min versus the ~ 13 hr projected for the original fly.io-routed orchestration.

Verdict: NO_CLOSURE. The gap drops 9% across the full $16\times$ refinement and then flatlines:

N_{steps}	mean gap	SEM of mean	max gap	PASS / 8
1000	0.1748	0.0091	0.2126	0/8
2000	0.1665	0.0083	0.2020	0/8
4000	0.1607	0.0089	0.2007	0/8
8000	0.1600	0.0087	0.1997	0/8
16000	0.1598	0.0090	0.2025	0/8

Every one of the 40 trajectories returns FAIL; the linear extrapolation at the asymptotic shrinkage rate would not reach the 0.02 tolerance until $N \sim 10^{10}$ steps. The trajectories themselves *do* converge: the Flyvbjerg–Petersen blocking regime transitions from `no_plateau` at $N = 1000$ to `plateau` at $N = 16000$, with n_{eff} growing from 7.8 to 125. The time-average is well-defined; it simply equals $\langle P \rangle_{\text{time}} \approx 0.66$, not $\langle P \rangle_{\text{heatbath}} \approx 0.506$.

Physical reading. Shell $\neq$ ensemble at the buckyball’s finite size. Each symplectic trajectory’s energy shell has a time-averaged plaquette ≈ 0.66 ; the canonical Boltzmann average over all energy shells weighted by their density of states gives ≈ 0.506 . The cross-check is not a check on whether the integrator is correct (it is — the eight v1.2 PASSes establish that); it is a probe of the gap between microcanonical and canonical

equilibrium at the buckyball’s $V = 60 / E = 90$ scale. That gap is ≈ 0.16 , stable under $16\times$ refinement of trajectory length, and is the open caveat the chapter carries.

Reproducibility.

```
python -m inertia_damping.scripts.section5_cuda_sweep
python -m inertia_damping.scripts.section5_closure_analysis \
    --sweep-root inertia_damping/reports/section5_sweep_cuda
```

The artefact of record is `inertia_damping/reports/section5_sweep_cuda/section5_closure.json` SHA-256 `145e331340f657e3eacf4704c9bb04a4181972b988e25b88114af9d98d623183`. The polished narrative is `SECTION5_CLOSURE_RECEIPT.md` alongside, with 40 per-trajectory sidecars (each carrying the full Section 5 schema) and two matplotlib figures (`section5_gap_vs_nsteps.pdf`, `section5_pt_vs_ph.pdf`).

A.7 The Falsification Battery (sudoku elimination grid)

Appendix A.6 derived the load-bearing experimental claim: at fixed static gravitational load, $\partial_Q \mu_Q \neq 0$ with K_Q , c_Q , drive amplitude, and systematics independently fit and controlled. By itself this is a positive assertion, not a falsifiable one. The Falsification Battery operationalises the Davis Method’s *Sudoku Principle* (Principle 5: elimination, not construction) by enumerating the null space and striking through what the simulation can rule out. The completion invariant — the count of struck rows in the simulatable subset — is what makes the claim load-bearing. The full specification is `inertia_damping/HALCYON_FALSIFICATION_BATTERY_SPEC.md` (v2), design-locked after a four-critic adversarial review that closed nine blocker-level defects in the v1 draft.

The grid. Ten null hypotheses $H_0 \dots H_9$ form the rows of the elimination grid. Seven are simulatable; three (thermal, EM, and mechanical pickup) require hardware DOFs the simulation does not have, so the simulation states the prediction-of-absence and the hardware apparatus is responsible for ruling them out at its own scale.

#	Null	Protocol	Gate threshold	Realm
H_0	nothing happens	drive sweep, multi- Q	$ \alpha > 6\sigma_\alpha$	sim
H_1	Eötvös material effect	vary μ_{proxy} on μ_{baseline}	$ d\alpha/d\mu_{\text{proxy}} / \alpha < 0.05$	sim
H_2	thermal pickup	—	—	hw only
H_3	EM pickup	—	—	hw only
H_4	mount/mechanical pickup	—	—	hw only
H_5	drive-amplitude artefact	amplitude scan, F -test	$F < F_{\text{crit}}(0.05)$	sim
H_6	single- ω resonance	full $\chi_Q(\omega)$ fit	$\chi^2/\text{dof} < 1.5$, dof= 21	sim
H_7	statistical fluctuation	FP blocked SEM	rel SEM $\leq 5\%$ of α	sim
H_8	Q sector drift	continuous $Q_{\text{surrogate}}$ monitor	$\sigma_Q/\langle Q \rangle < 0.03$	sim
H_9	τ_Q model error	alternative functional form	$ \alpha_{\text{alt}} - \alpha / \alpha < 0.20$	sim

The substrate model the battery probes. The test-mass dynamics is governed by $\mu_{\text{total}}(Q) \ddot{x} + c_Q \dot{x} + K_Q x = F(t)$, where $\mu_{\text{total}}(Q) = \mu_{\text{proxy}} \cdot \mu_{\text{baseline}} + \mu_{\text{eff}}(Q)$. The Halcyon coupling enters only through $\mu_{\text{eff}}(Q)$, computed on the buckyball substrate via

$$\mu_{\text{eff}}(Q) = \alpha_{\text{Halcyon}} \frac{1}{E} \sum_{e=1}^{90} \kappa_Q(e) \tau_Q^2(e) |\phi_n(e)|^2, \quad (1)$$

with $\kappa_Q(e) = (1 - q_0(U_{f_1(e)}))(1 - q_0(U_{f_2(e)}))$, $\tau_Q(e) = \tau_0/(1 + \beta_\tau s_Q(e))$, and $\phi_n(e)$ the n -th non-trivial eigenmode of the buckyball graph Laplacian projected to edges via gradient. The model is material-independent by construction: μ_{proxy} modulates the bare baseline only, so a non-zero $\partial_Q \mu_{\text{total}}/\partial \mu_{\text{proxy}}$ would indicate that what the simulation measures is a material artefact, not Halcyon. This is the H_1 test, and it discriminates because the model is designed to fail it cleanly if μ_{proxy} were allowed to enter μ_{eff} (an earlier draft of the SPEC made that mistake; the closure of v1's H_1 blocker is the load-bearing fix that earned the v2 design lock).

The simulation kernel is `inertia_damping/test_mass_dynamics.py`, a new module (not an extension of the Section 5 CUDA kernel, which is pure-gauge and supports no external DOFs) that couples the buckyball SU(2) gauge leapfrog to a single-DOF driven damped oscillator via the $\mu_{\text{eff}}(U)$ recomputed at every step. The vectorised face holonomies from the Section 5 CUDA kernel are reused on CPU with $B = 1$, giving a $12\times$ speedup over the per-step Python loop and bringing the per-lock-in wall to ~ 40 s. The orchestrator `inertia_damping/falsification.battery.py` runs the sweep design (5 frequencies $\times$ 3 Q sectors $\times$ 8 seeds for the main grid; additional amplitude scan for H_5 , μ_{proxy} scan for H_1 , alternative- τ_Q scan for H_9) and emits the `section_11_falsification.battery` JSON schema specified in SPEC v2 §6.

Smoke results (8 seeds, gauge enabled, run `battery_fast_20260620_104846`, 78.5 min wall):

- Sector separation: PASS.
- $\alpha_{\text{measured}} = 5.92 \times 10^{-5} \pm 3.8 \times 10^{-6}$, $|\alpha|/\sigma_\alpha = 15.62$.
- Per-gate verdicts:
 - H_0 (nothing happens): STRIKE ($|\alpha|/\sigma_\alpha = 15.62 \gg 6$).
 - H_1 (material): STRIKE ($|d\alpha/d\mu_{\text{proxy}}|/|\alpha| = 0.028 < 0.05$).
 - $H_2/H_3/H_4$: N/A (hardware-only realm).
 - H_5/H_6 : N/A (not run in `--battery-fast`; full sweep is open).
 - H_7 (statistics): NOT STRUCK (rel. α -SEM 6.4%, vs. 5% gate).
 - H_8 (Q drift): STRIKE (mean rel. drift 1.7%, max 4.3%, vs. 3% gate; mean is well under).
 - H_9 (τ_Q model error): STRIKE ($|\alpha_{\text{alt}} - \alpha|/|\alpha| = 0.025$, well under 20% gate; tightened from the 3-seed run's 0.076).
- Sudoku verdict: FAIL_NULL_SURVIVES. 4-of-5 simulatable nulls struck; H_7 remains marginal.

Why H_7 did not close at 8 seeds (an honest diagnostic, not a paper over). Going from 3 to 8 seeds moved the rel. α -SEM only from 6.6% to 6.4%. Pure seed-replication noise should have shrunk like $1/\sqrt{N_{\text{seed}}}$, i.e. by a factor of $\sqrt{8/3} \approx 1.6$, not stayed constant. The dispersion floor is therefore *not seed-dominated*; it comes from the current α -extraction averaging over multiple drive frequencies ω as if they were seed replicates. Each frequency probes a different point on the $\chi_Q(\omega)$ curve, and the implied $d\mu_{\text{eff}}/dQ$ slope at fixed ω differs slightly between the quasi-static ($\omega = 0.3\omega_c$), inertial ($\omega = \omega_c$), and impulsive ($\omega = 3\omega_c$) regimes. The clean fix is to extract $\alpha(\omega)$ per-frequency first,

then combine; that becomes the H_6 full- $\chi(\omega)$ -fit protocol the full battery runs, which is structurally what closes the gate. The marginal H_7 is therefore a faithful signal that the `--battery-fast` mode bumps into the limit of its own extractor, not that the signal is statistical noise — the actual signal-to-noise on α is $|\alpha|/\sigma = 15.62$. The chapter’s verdict is honest about this: the simulation passes 4-of-5 simulatable nulls cleanly, and the fifth’s marginal failure is the extractor’s diagnostic, not a substrate failure.

Self-consistency vs. independent prediction. The simulation can compute $\alpha_{\text{predicted}}^{\text{self-consistency}} = \partial\mu_{\text{eff}}/\partial Q|_{\text{default}}$ from finite differences over the three Q labels, and this self-consistency value matches α_{measured} trivially (both are derived from the same model). *It is therefore not a load-bearing falsifier.* The independent $\alpha_{\text{predicted}}^{\text{independent}}$ — a closed-form derivation from the Davis Field Equations evaluated at $\beta = 2.5$ on the buckyball substrate — is open work, tracked as the next item in the chapter’s open queue, and is what the eventual hardware bench will compare its α_{measured} against. The simulation closes the simulation-side gates only; it cannot close the independent-prediction gate.

The ergodicity caveat. Per A.5, the buckyball substrate exhibits a documented $\sim 16\%$ irreducible shell-vs-ensemble gap. The battery operates within the microcanonical energy shell sampled by the seed-and-canonical-init protocol; a `PASS_SIMULATION_ONLY` verdict therefore does not imply thermodynamic equilibrium. For hardware extension to a larger lattice the gap is expected to vanish under the law of large numbers and the battery must re-validate at the new scale; the SPEC does not certify cross-scale invariance.

The receipt. The 8-seed `--battery-fast` sidecar is `inertia_damping/reports/section11_battery` SHA-256 `b841ede05cffabfffb82efc1b21b63ed9e22341125b01e3f16801d4412c789c`.

A.7.1 The full-battery follow-up: an honest falsifier of the extractor itself

The smoke results above used a slope-through-origin extractor on the internal `mu_eff_mean` value tracked by the simulation. That number is the model’s own internal state, not the value the apparatus would observe. SPEC v2 §4 specifies a stricter protocol — the H_6 test: fit $\chi_Q(\omega) = [K_Q + ic_Q\omega - \mu_Q\omega^2]^{-1}$ per Q to the lock-in-demodulated $|\chi|$ and $\arg \chi$ curves at multiple ω , then extract α from the fitted μ_Q values. The fitted μ_Q is what a real experiment would see; the internal $\mu_{\text{eff}}(U(t))$ is what only the simulator knows. The two extractors are different epistemic objects, and the SPEC’s choice to make the observation-space extractor load-bearing was the right one.

A full-battery run with the H_6 extractor wired up (sidecar `battery_full_20260620_181227.json`, SHA-256 `ad33fd568f6257ff0ba98fd9d31a22403cb9915de083a6b6ca1384c49e4b49b8`, 111.2 min wall, $5 \omega \times 3 Q \times 8$ seeds main grid plus $H_1/H_5/H_9$ supplements) returned:

- Sector separation: `PASS` (17.83σ minimum pairwise).
- Per- Q measured $\chi(\omega)$ curves are *identical to four decimal places across Q* : at $\omega = \omega_c$, $|\chi|_{Q=0} = 6.0519$, $|\chi|_{Q=1} = 6.0519$, $|\chi|_{Q=2} = 6.0522$. The Halcyon-predicted shift at the SPEC’s default coupling $\alpha_{\text{Halcyon}} = 1.0$ (lattice units) is roughly 60 ppm of the baseline susceptibility — *below the measurement noise floor of the observation-space extractor*.
- Sudoku verdict: `FAIL_SIGNAL_MISSING`. H_0 does not strike: $|\alpha_{\text{fitted}}|/\sigma_\alpha = 0.21$, dominated by fit noise. The chapter does not claim a signal here.

- H_1 (material): STRIKE (rel slope 0.0046; geometric independence robust under this extractor too).
- H_5 (amplitude linearity): STRIKE (rel slope χ vs F_0 effectively 0).
- H_6 ($\chi(\omega)$ fit quality): NOT STRUCK ($\chi^2/\text{dof} = 1.3 \times 10^8$; the fit explodes because the curves cannot discriminate three Q values to within their dispersion).
- H_7 (statistics): NOT STRUCK, rel α -SEM 474%.
- H_8 (Q drift): NOT STRUCK, mean drift 18.5%, max 100%. The initial Q -bias relaxes back toward the trivial vacuum during the longer measurement window.
- H_9 (τ_Q model robustness): STRIKE (rel diff 12.6%, well under 20% gate).

The interpretation. The smoke-mode result was *internally* consistent but *externally* unmeasurable at the SPEC’s default coupling. The two extractors give different verdicts because they ask different questions:

- `extract_alpha` (slope through origin on internal μ_{eff}): “Does the model’s own state respond to Q ?” — answer YES at $|\alpha|/\sigma = 15.62$.
- `extract_alpha_from_fits` (per- Q $\chi(\omega)$ fit on observed susceptibility): “Would a lock-in measuring the test mass see the response?” — answer NO at $|\alpha|/\sigma = 0.21$.

Both are correctly implementing what they claim to measure. The SPEC’s choice to make the observation-space extractor load-bearing is what surfaces the calibration gap: at $\alpha_{\text{Halcyon}} = 1.0$, the predicted $\mu_{\text{eff}}/\mu_{\text{baseline}}$ ratio at $Q = 2$ is $\sim 6 \times 10^{-5}$ — too small to climb above the lock-in noise floor of a ~ 30 time-unit driven trajectory.

This is the actual scientific content of A.7.1: *the battery is doing its job by failing*. A real signal needs either (a) a stronger predicted coupling — α_{Halcyon} calibrated from the Davis Field Equations rather than set to 1.0 by convention; (b) a longer measurement window that the gauge dynamics do not relax through (the H_8 drift); or (c) an actively driven Q sector that the apparatus pins, not the passive initial-condition bias the simulation currently uses. All three are open work tracked outside this chapter. The simulation cannot supply the independent value of α_{Halcyon} ; that comes from the Field Equations.

What this commits the chapter to. The simulation has demonstrated that the SPEC v2 extractor discipline is operative: it refuses to score a passing verdict when the signal is below the observation floor. The four-of-five strikes in `--battery-fast` reflect internal model consistency (which is real and useful for verifying the framework’s mechanics), not an observable apparatus claim. The three-of-seven strikes in the full battery reflect the apparatus-side reality at the current calibration. Both are honest readings of what each test asks. The chapter’s load-bearing claim — $\partial_Q \mu_Q \neq 0$ in observation space — is not yet substantiated by this simulation at the SPEC’s default α_{Halcyon} and will require either re-calibration or a longer-window measurement protocol. That is the open item the next chapter version will address.

A.7.2 The calibrated run: $\alpha_{\text{Halcyon}} = 1000$ surfaces the second blocker

A.7.1’s diagnosis pointed at $\alpha_{\text{Halcyon}} = 1.0$ as a calibration too small to produce an observable shift. A frozen-gauge calibration scan confirmed that the per- Q instantaneous

$|\chi|$ spread scales monotonically with α : $\alpha = 1$ gives a spread of 0.005 (under noise), $\alpha = 100$ gives 0.41, $\alpha = 1000$ gives 1.60, $\alpha = 10000$ saturates near the unphysical limit. The natural follow-up is the second full-battery run at $\alpha_{\text{Halcyon}} = 1000$, on the contracted ω grid $\{1.0, 3.0, 10.0\} \times \omega_c$ (dropping the quasi-static cells where the period exceeded t_{total}).

Sidecar of record: `battery_calibrated_20260621_011304.json`, SHA-256 `8183edfa936aeddffa3ba719ce6971f1237e3f0108bd5eba70e59dd558575d037`, 55.5 min wall, 8 seeds, 3 ω , 3 Q , $n_{\text{steps}} = 1200$, $\alpha_{\text{Halcyon}} = 1000$, gauge dynamics live. The result:

- Verdict: FAIL_SIGNAL_MISSING again. $|\alpha_{\text{fitted}}|/\sigma_\alpha = 0.52$ — only $2.5\times$ better than the $\alpha = 1$ run despite the $1000\times$ bump in coupling.
- H_5 (amplitude linearity): STRIKE (rel slope effectively zero).
- H_6 ($\chi(\omega)$ fit quality): NOT STRUCK, χ^2/dof dropped 5 orders of magnitude (from 1.3×10^8 to 1777). The fit is no longer hopeless — the per- Q curves are now distinguishable in principle — but they are still not consistent with the 3-parameter Lorentzian.
- H_8 (Q drift): NOT STRUCK. Mean rel. drift 20.5%, max 128%. Not improved by the α bump — this is a separate physics constraint.
- H_1, H_9 : strike as before, by the internal extractor.

The diagnosis sharpens. The α bump did exactly what a calibration bump should do: it moved the predicted signal up by $1000\times$. But the observation-space SNR went up only $2.5\times$, not $1000\times$. The difference — a factor of ~ 400 — is the noise floor’s own scaling with α . The same gauge-field dynamics that produce $\mu_{\text{eff}}(Q)$ also produce the H_8 drift; bumping α amplifies both. The signal-to-noise ratio asymptotes to the timescale ratio

$$\frac{S}{N} \sim \frac{\tau_{Q\text{relaxation}}}{t_{\text{total}}},$$

not to α . Pushing α harder will not close H_8 — only changing the measurement-window-to-relaxation-time ratio will.

The real blocker, named cleanly. The simulation as designed has no active Q -pinning mechanism. The buckyball Hamiltonian’s minimum-energy state is the trivial vacuum at $Q_{\text{surrogate}} = 0$, and any non-trivial initial condition (the simulation’s $Q = 1, 2$ sectors) is a perturbation that relaxes back during measurement. A real apparatus pins Q through the cage’s continuous drive of the Josephson junctions — the cage doesn’t *prepare* Q and then *watch* it decay; it *holds* Q throughout measurement. The simulation’s passive initial-condition protocol is the wrong model for this.

Three paths forward, named for SPEC v3.

1. **Add an active Q -restoring potential** $V_{\text{pin}}(U) = \lambda_{\text{pin}} (Q_{\text{surrogate}}(U) - Q_{\text{target}})^2$ to the gauge Hamiltonian. This models the cage’s continuous drive. Changes the dynamics; needs unit tests to verify symplecticity and to confirm the pinning does not itself contaminate the lock-in via parametric coupling.
2. **Use a measurement window shorter than the Q -relaxation timescale**, with proportionally higher drive frequencies. Requires re-design of the sweep grid: ω -min

set by “at least 5 cycles in t_{short} ,” t_{short} set by “ Q drift $< 3\%$ over the window.” Likely $t_{\text{short}} \approx 3$ time units, $\omega \geq 10$, drive cycles in $\{2, 5, 10\} \times \omega_c$. Trade-off: fewer cycles per cell means noisier demodulation per cell, so more seeds needed.

3. **Accept H_8 as a hardware-only gate**, in the same spirit as $H_2/H_3/H_4$. The simulation predicts the absence of Q -pinning is the dominant blocker; the apparatus is responsible for closing H_8 through its drive design.

This chapter does not pick a winner among the three. SPEC v3 is the next design-doc revision; the simulation-side measurement architecture is the next implementation revision. The chapter v4.4 ships with A.7.1 and A.7.2 as the honest 3-stage diagnostic: smoke mode read internal state (passed), full mode with the proper extractor caught the calibration gap (revealed α -scale was unobservable), and the calibrated run with the bumped α caught the H_8 drift as the deeper blocker (revealed the simulation’s passive- Q -bias protocol is the wrong model for the apparatus). **Each stage of the battery, by failing, makes the experimental claim more falsifiable, not less.** The SPEC v2 design document is `inertia_damping/HALCYON.FALSIFICATION.BATTERY.SPEC.md`. Reproducibility command:

```
python -m inertia_damping.scripts.run_battery --fast
```

The standalone script `inertia_damping/scripts/run_battery.py` ships in this commit; integration into `run_validation_report.py` as a first-class `--battery` flag is the next item in the open queue.

A.6 The thread experiment and the dynamic transfer-function target

The textbook two-thread inertia demonstration [1, Vol. I, §10–2] is the simplest geometry that already distinguishes the two readings of μ Halcyon needs to separate: μ -as-a-property versus μ -as-a-coupling. A weight hangs from a thread tied to the ceiling; a second thread hangs from the bottom of the weight to a hand. Pull the lower thread slowly and the upper thread breaks (it bears $mg + F$ in static equilibrium). Pull the lower thread fast and the lower thread breaks (the mass cannot redistribute the impulse fast enough; the lower thread bears F alone). Same weight, same gravity, different breaking thread. The control parameter is the Deborah number [2]

$$\mathcal{D} = \frac{\tau}{t_{\text{load}}} = \frac{r \tau}{F_*} \quad (2)$$

where $r = dF/dt$ is the loading rate, $t_{\text{load}} = F_*/r$, F_* is the breaking-force scale, and τ is the bulk’s inertial relaxation time. The thread experiment lives in the two limits $\mathcal{D} \ll 1$ (quasi-static, upper bears $mg + F$) and $\mathcal{D} \gg 1$ (impulsive, lower bears F alone).

The idealised inertial crossover. In an elastic model with no dissipation, the crossover from quasi-static to inertial response is set by the balance $K \sim \mu \omega^2$, giving

$$\omega_c = \sqrt{K/\mu}, \quad \tau_\mu = \sqrt{\mu/K}, \quad \mu = K \tau_\mu^2. \quad (3)$$

The dimensional discipline matters: τ_μ is the *inertial* relaxation time (the time for an impulse to propagate through the bulk), not the viscous relaxation time $\tau_\eta = \eta/K$.

Conflating the two is a common error; the thread demo, having no viscous element in the idealisation, isolates τ_μ cleanly.

The real apparatus. For finite damping, the linear-response (Kubo) susceptibility for the single relevant degree of freedom [3, §6.4] is

$$\chi_Q(\omega) = \frac{1}{K_Q + i c_Q \omega - \mu_Q \omega^2}. \quad (4)$$

For a measured normal mode of shape $\phi_n(x)$, the mode-effective inertial coefficient is

$$\mu_{\text{eff}}^{(n)}(Q) \sim \int \kappa_Q(x) \tau_Q^2(x) |\phi_n(x)|^2 dV. \quad (5)$$

The $|\phi_n|^2$ projection makes the coordinate dependence explicit; effective mass is mode-dependent in any extended elastic system.

The falsifiable claim. Halcyon’s hypothesis, written cleanly:

$$\partial_Q \mu_Q \stackrel{?}{\neq} 0 \quad (6)$$

at fixed K_Q , c_Q , drive amplitude, thermal drift, and EM systematics. Three implementation discipline points:

1. *Fit the full transfer function.* A shift in $\omega_c(Q) = \sqrt{K_Q/\mu_Q}$ alone proves only a shift in the ratio K_Q/μ_Q . The clean inertial-shift claim requires fitting K_Q , μ_Q , c_Q as independent parameters.
2. *Include damping.* The idealised model can omit c_Q ; the apparatus model cannot. Mechanical mounts, eddy currents, and air loading all contribute finite damping that must be either physically suppressed or fit and subtracted.
3. *Preregister systematics.* Drive amplitude, thermal drift, magnetic pickup, vibration, and cage drift constitute the systematics budget. Eq. 6 is a positive claim only after that budget is set in advance and the residual slope exceeds it. This is the same discipline that resolved MICROSCOPE [4].

The equivalence-principle framing. This is an equivalence-principle-sensitive measurement, not an equivalence-principle-preserving one. MICROSCOPE bounds m_i/m_g at the 10^{-15} level between materials at fixed material state [4]; Halcyon tests whether μ_Q depends on the programmed gauge sector of the surrounding apparatus at fixed material and fixed static gravitational load. A non-zero $\partial_Q \mu_Q$ is an EP-sensitive result — not a contradiction-free preservation of $m_i = m_g$. Halcyon is by construction a *state-programmed equivalence-principle stress test*: same test mass, same gravitational load, different programmed gauge sector.

The distinction from prior claims. The historical class Halcyon distinguishes itself from — the Woodward effect, Podkletnov’s rotating disc, the EmDrive — has consistently failed under careful measurement; the records are consistent with zero modulo systematics. Halcyon’s response is not a stronger version of the same target. The target is different: a dynamic transfer-function shift in Eq. 6, not a static weight or thrust anomaly. The scale never notices. A weight measurement is the wrong instrument; a lock-in driven response measurement at multiple frequencies, fitted across multiple Q , is the right one.

The connection to the buckyball substrate. The Halcyon substrate enters Eq. 5 through $\kappa_Q(x)$ and $\tau_Q(x)$. The buckyball cage realises the programmed gauge sector Q as a Wilson-action ground state on the $V = 60$, $E = 90$, $F = 32$, $\chi = 2$ truncated icosahedron, $SU(2)$, $\beta = 2.5$ substrate this chapter validates. The cage geometry sets κ_Q ; the local field configuration sets τ_Q . The falsification target Eq. 6 is the only experimental statement that depends on the Halcyon framework being right rather than on standard lattice-gauge bookkeeping. Everything else — the Migdal–Witten target, the eight PASS verdicts in §3, the Section 5 closure receipt in A.5 — is preparation. A.6 is the actual measurement bridge.

References

- [1] R. P. Feynman, R. B. Leighton, and M. Sands, *The Feynman Lectures on Physics*, Vol. I (Addison–Wesley, 1963), §10–2.
- [2] M. Reiner, “The Deborah Number,” *Physics Today* **17**(1), 62 (1964).
doi:10.1063/1.3051374.
- [3] H. Goldstein, C. P. Poole, and J. L. Safko, *Classical Mechanics*, 3rd ed. (Addison–Wesley, 2002), §6.4.
- [4] P. Touboul *et al.* (MICROSCOPE collaboration), “MICROSCOPE Mission: Final Results of the Test of the Equivalence Principle,” *Phys. Rev. Lett.* **129**, 121102 (2022).
doi:10.1103/PhysRevLett.129.121102.